

# Amazon EC2 + EWS Lambda

---

## Slide 1: Where Can You Store Data?

### Slide Content (as text):

- **Instance Store:** Temporary, high-speed storage tied to the instance lifecycle
- **EFS (Elastic File System)** Support: Shared file storage
- **EBS (Elastic Block Storage):** Persistent block-level storage
- **S3:** Large data set storage or EC2 backups

### Expanded Explanation:

1. **Instance Store**
  - This is storage physically attached to the host machine running your EC2 instance.
  - It is extremely fast and ideal for temporary or ephemeral data.
  - **Important Note:** If you stop or terminate the EC2 instance, data on the instance store is lost.
2. **EFS (Elastic File System)**
  - A fully managed, shared file system that can be mounted on multiple EC2 instances simultaneously.
  - Ideal for scenarios where you need a common file system for multiple servers or services.
3. **EBS (Elastic Block Storage)**
  - Persistent storage volumes that can be attached to an EC2 instance.
  - Data remains available even if you stop or reboot the instance (but not if you terminate and delete the volume, unless you configure otherwise).
  - Great for typical server use-cases like a root volume or attached volumes for databases.
4. **S3 (Simple Storage Service)**
  - Object-based storage for large amounts of data (files, images, backups, static web content, etc.).
  - Highly durable and cost-effective.
  - Typically used for backups, hosting static sites, or data lakes.

---

## Slide 2 (labeled as "6" in the screenshot): Amazon EC2 & Lambda

### Slide Content (as text):

- **DS 4300**
- **Amazon EC2 & Lambda**
- **Mark Fontenot, PhD**
- **Northeastern University**

**Expanded Explanation:**

- This slide introduces the topic of Amazon EC2 (Elastic Compute Cloud) and AWS Lambda.
  - **EC2** is Amazon's Infrastructure as a Service (IaaS) offering, allowing you to rent virtual machines ("instances") in the cloud.
  - **AWS Lambda** is a serverless computing service that lets you run code without provisioning or managing servers.
- 

**Slide 3 (labeled "7" in the screenshot): Let's Spin Up an EC2 Instance**

**Slide Content (as text):**

- "Let's Spin Up an EC2 Instance" (title)

**Expanded Explanation:**

- This slide indicates the beginning of a tutorial or demonstration on how to create ("spin up") a new EC2 instance.
  - Spinning up an instance involves choosing an AMI (Amazon Machine Image), selecting instance type, configuring network settings, etc.
  - The upcoming slides likely detail the step-by-step process.
- 

**Slide 4 (labeled "8"): Let's Spin Up an EC2 Instance**

**Slide Content (as text, partially visible):**

- Screenshots from the AWS Console showing steps to create a new EC2 instance.

**Expanded Explanation:**

- **Step 1: Choose an AMI** – For example, "Ubuntu Server 20.04 LTS".
- **Step 2: Choose an Instance Type** – e.g., t2.micro for a small, free-tier-eligible instance.
- **Step 3: Configure Instance Details** – VPC settings, subnets, etc.
- **Step 4: Add Storage** – EBS volume size, etc.
- **Step 5: Add Tags** – Optional metadata for better organization.
- **Step 6: Configure Security Group** – Control inbound/outbound traffic with firewall rules.

- **Step 7: Review & Launch** – Launch the instance and create or select a key pair for SSH access.
- 

## Slide 5 (labeled "9"): Let's Spin Up an EC2 Instance

### Slide Content (as text):

- Another screenshot continuing the setup steps.

### Expanded Explanation:

- Likely continues detailing the creation process, focusing on security group configuration or final review.
  - **Security Group** best practice:
    - Open only the ports you need.
    - Typically, port 22 for SSH (for Linux) or RDP (for Windows) is required.
    - If you plan to run a web server, open port 80 (HTTP) or 443 (HTTPS).
- 

## Slide 6 (labeled "10"): Ubuntu VM Commands

### Slide Content (as text):

- **Initial user is ubuntu**
- Access super user commands with sudo
- Make sure apt-get is up to date
- Update the packages installed
- sudo apt update; sudo apt upgrade

### Expanded Explanation:

- Once you SSH into the instance, you'll be logged in as the ubuntu user by default (on an Ubuntu AMI).
  - To perform administrative tasks (install packages, change configurations), you'll use sudo.
  - **Key commands:**
    - sudo apt update – Refreshes the package lists for upgrades/new packages.
    - sudo apt upgrade – Installs available updates.
  - This ensures your instance is secure and has the latest security patches.
-

## Slide 7 (labeled "11"): MiniConda on EC2

### Slide Content (as text):

- Make sure you're logged in to your EC2 instance.
- Let's install MiniConda:
  - Example command: `wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh`
  - Then run `bash Miniconda3-latest-Linux-x86_64.sh`.

### Expanded Explanation:

- **MiniConda** is a minimal installer for Conda, a package manager that helps you install Python and data-science-related libraries easily.
  - After installation, you typically need to source your `~/.bashrc` or restart your shell so that the conda command is recognized.
  - You can then create or manage virtual environments on your EC2 instance.
- 

## Slide 8 (labeled "12"): Installing & Using Streamlit

### Slide Content (as text):

- Log out of your EC2 instance and log back in.
- streamlit is now available!
- In your environment, install Streamlit.
- Make a test Streamlit file.
- Basic usage: `streamlit run my_app.py`
- `cd web`

### Expanded Explanation:

- After installing Streamlit with `pip` or `conda install streamlit`, you can run it by calling `streamlit run <filename.py>`.
  - You might log out and back in to refresh your `PATH` or environment variables.
  - Creating a quick test file (e.g., `my_app.py`) helps confirm that Streamlit is installed correctly.
- 

## Slide 9 (labeled "13"): Basic Streamlit App

### Slide Content (as text):

```
import streamlit as st
```

```
st.title("My Streamlit App")  
st.write("Hello, world!")
```

- Save as test.py
- nano test.py (to edit)
- streamlit run test.py

#### **Expanded Explanation:**

- **Streamlit** is a Python library for building simple web apps quickly, especially for data visualization and data science demos.
  - A minimal example includes a title and a simple "Hello, world!" text.
  - st.title() sets the main heading, while st.write() can display text, data frames, or other content.
- 

### **Slide 10 (labeled "14"): Opening Up The Streamlit Port**

#### **Slide Content (as text):**

- Screenshot of AWS security group settings.
- Possibly showing how to open the port used by Streamlit (default is 8501).

#### **Expanded Explanation:**

- By default, Streamlit runs on localhost:8501.
  - To make it accessible externally, you can either:
    1. Use an SSH tunnel (not covered here), or
    2. Open port 8501 in your EC2 security group inbound rules so you can access http://<public-ip>:8501.
  - **Security Group:**
    - Add a rule:
      - Type: Custom TCP
      - Port Range: 8501
      - Source: Your IP (or 0.0.0.0/0 if you want it accessible to everyone—less secure)
- 

### **Slide 11 (labeled "15"): In a Browser**

#### **Slide Content (as text):**

- Screenshot of the Streamlit app running in the browser, with a title "Welcome to my Streamlit App."

**Expanded Explanation:**

- After you've opened the port and started Streamlit on the server (streamlit run test.py), navigate to `http://<EC2-Public-IP>:8501`.
  - You should see your basic Streamlit app.
  - This confirms your setup is correct: instance is running, port is open, and the app is served.
- 

**Slide 12 (labeled "16"): AWS Lambda**

*(This slide seems to be a title or transition slide.)*

**Slide Content (as text):**

- "AWS Lambda" (title)

**Expanded Explanation:**

- Now shifting the topic from EC2 to AWS Lambda.
  - Lambda is a *serverless* compute service where you only pay for the time your code is running.
  - You don't need to manage any infrastructure; AWS handles the underlying servers.
- 

**Slide 13 (labeled "17"): Lambdas**

**Slide Content (as text):**

- Lambdas provide **serverless computing**.
- Automatically run code in response to events.
- Worry about the code, not the servers.
- You only pay for execution time, not idle compute time (different from EC2).

**Expanded Explanation:**

- **Serverless** means you don't provision or manage servers. Instead, you write functions that AWS executes on-demand.
- **Events** can be triggered by S3 uploads, API Gateway calls, DynamoDB updates, scheduled CloudWatch events, etc.
- **Pay-as-you-go** model:
  - Billed by the number of requests and the time your code runs.

- No charges when your code isn't running.
- 

## Slide 14 (labeled "18"): Lambda Features

### Slide Content (as text):

- **Event-driven execution** – can be triggered by many different events in AWS.
- Supports a large number of runtimes: Python, Java, Node.js, etc.
- **Highly integrated** with other AWS services.
- **Extremely scalable** and scales automatically.

### Expanded Explanation:

- Because Lambda is event-driven, you don't have to keep an instance running.
  - Official runtimes include Node.js, Python, Ruby, Java, Go, .NET, etc. You can also use custom runtimes.
  - Integrations with S3, SNS, SQS, DynamoDB, API Gateway, and more let you build sophisticated serverless architectures.
  - Lambda automatically handles concurrency. If multiple events occur at once, Lambda will spin up additional executions.
- 

## Slide 15 (labeled "19"): How it Works

### Slide Content (as text):

1. Add/Upload your code through AWS Management Console (or via CLI, or CI/CD).
2. Configure event source(s).
3. Watch your Lambda run when one of the event sources triggers it.

### Expanded Explanation:

- **Step 1:** You can write your function code in the console's built-in editor, upload a ZIP, or reference an S3 bucket.
  - **Step 2:** An event source (e.g., S3 bucket event, CloudWatch event) is configured to call your function.
  - **Step 3:** Once the event occurs, AWS Lambda automatically executes your function.
- 

## Slide 16 (labeled "20"): Let's Make One

*(Likely an introduction to a hands-on Lambda creation demo.)*

**Slide Content (as text):**

- "Let's Make One" (title)

**Expanded Explanation:**

- You will walk through creating a new Lambda function in the AWS Console.
  - The following slides presumably show each step in the process.
- 

**Slide 17 (labeled "21"): Making a Lambda**

**Slide Content (as text):**

- Screenshot from the AWS Lambda console, showing the "Create function" button or the initial steps in creation.

**Expanded Explanation:**

- On the Lambda main page, click **Create function**.
  - You'll be prompted to choose between:
    1. **Author from scratch**
    2. **Use a blueprint**
    3. **Use container image**
    4. **Browse serverless app repository**
  - Typically, you'll "Author from scratch" to get started with a basic function.
- 

**Slide 18 (labeled "22"): Creating a Function**

**Slide Content (as text):**

- Possibly shows form fields:
  - Function name
  - Runtime (Python, Node.js, etc.)
  - Permissions/Role

**Expanded Explanation:**

- **Function name:** Choose something descriptive, e.g., my-first-lambda.
- **Runtime:** Select the language environment (e.g., Python 3.9).
- **Permissions:** You can create or select an IAM role that gives Lambda the permissions it needs (for example, to log to CloudWatch).



---

## Slide 19 (labeled "23"): Sample Code

### Slide Content (as text):

```
def lambda_handler(event, context):  
    # Your code here  
    return {  
        'statusCode': 200,  
        'body': 'Hello from Lambda!'  
    }
```

### Expanded Explanation:

- This is a simple Python Lambda handler.
- The `lambda_handler` function is the entry point for the Lambda.
- **Parameters:**
  - `event`: data passed in by the triggering service or event.
  - `context`: runtime information about the function execution (e.g., request ID, memory limits, etc.).
- The returned dictionary typically includes a `statusCode` and `body` (especially if used behind an API Gateway).

---

## Slide 20 (labeled "24"): Edit the code, Deploy the code!

### Slide Content (as text):

- A screenshot showing the Lambda code editor in the console, with a "Deploy" button.

### Expanded Explanation:

- After writing or modifying your code in the console, you must **Deploy** it to save the changes and make them live.
- You can test the code by configuring a test event or by hooking up an actual event source.

---

## Slide 21 (labeled "25"): Test It

### Slide Content (as text):

- A screenshot showing the "Test" button and a sample event

configuration.

### Expanded Explanation:

- You can create a **test event** in the console that mimics the payload from a real event source (e.g., S3, API Gateway).
- Once you click **Test**, Lambda will run your function with the provided event.
- The console will display the result, including logs and the return value.

---

### Final Notes and Summary

By following these slides and the expanded explanations, you have a roadmap for:

1. **Creating and managing an EC2 instance:**
  - Choosing an AMI (like Ubuntu).
  - Updating and installing necessary packages (including MiniConda).
  - Deploying a basic web application (e.g., Streamlit).
  - Opening ports in the security group for external access.
2. **Introduction to AWS Lambda:**
  - Understanding serverless concepts.
  - Creating a basic Lambda function in Python.
  - Event-driven execution and the pay-per-execution model.
  - Testing your Lambda function with custom events.

These are foundational AWS concepts—EC2 for on-demand virtual machine instances and Lambda for serverless functions. With these, you can build a wide range of applications in the AWS ecosystem.

---

### End of Expanded Lecture Notes